

Abliteration Is Not a Scalpel: Off-Target Effects of Refusal Removal on Decision Disposition Across Model Families

Aleksander Fafuła, PhD — independent researcher, funduszai.pl Contact: aleksander@fafula.com

Draft v0.5 — 2026-07-19. Status: complete results for two model families (21,600 decisions); preregistered (prompt hash 305718d8..., PREREGISTRATION-v2.md, tag freeze-v2-2026-07-05). Numbers are final for the 18-week window 2026-03-07... 2026-07-04, all 18 weeks economically settled (the last Saturday settled on 2026-07-10).

Abstract

Abliteration — deleting a model's refusal direction from its weights — is the standard recipe behind popular "uncensored" open-weight models. It is advertised as surgical: remove refusal, change nothing else. We show the surgery is not clean. As a disposition probe we use 21,600 decisions under uncertainty — weekly up/down calls on 60 Warsaw Stock Exchange equities over 18 weeks, replayed through a frozen multi-agent pipeline so the decision-layer model is the only variable (5 samples per cell, zero generation failures). The task elicits no refusals at all — base arms completed all 10,800 of their decisions without refusing once — so there was no refusal behavior to remove, and any between-arm delta is pure side effect. Holding provenance constant — official BF16 checkpoints, a single abliteration author (huihui-ai), an identical serving stack, one byte-identical frozen prompt — we compare base and abliterated arms of two Mixture-of-Experts families (Gemma-4-26B-A4B-it and Qwen3-30B-A3B-Instruct-2507). Three effects replicate across both families (weeks-clustered bootstrap CIs excluding zero): abliterated models are systematically **more optimistic** (+12.2 pp Gemma, +7.4 pp Qwen more up-bets; the confirmed preregistered endpoint), **justify themselves at greater length**, and use **fewer explicit uncertainty words** with more concessive rhetoric in their self-critiques (both exploratory; length was preregistered as a covariate, not a directional hypothesis). A fourth effect **reverses sign**: the same operation makes Gemma-abliterated *less* confident and Qwen-abliterated *more* (family CIs non-overlapping) — one weight surgery, opposite shifts in expressed confidence. A capability-covariate analysis rules out instruction-following degradation as the driver, and no arm shows economic skill — the apparent edge of abliterated arms is regime beta, not alpha. Along the way our provenance audit caught two independent contamination channels (a mismatched-quantizer pilot pair, whose headline finding the clean design overturns, and a stale chat template in a community checkpoint that silently mangled the rendered prompt) — evidence that in studies of community-modified checkpoints, toolchain artifacts are the rule, not the exception. Abliteration's off-target footprint is real, partially universal, and model-dependent in direction: whoever deploys an "uncensored" model as an agent is deploying a measurably different decision-maker, not the base model minus refusals.

Key findings

1. **Removing refusals adds optimism.** On identical evidence, ablated models bet on the upside +7.4 to +12.2 pp more often — in both families and on both market segments (preregistered P1).
 2. **Same surgery, opposite confidence shifts.** Ablation makes Gemma *less* confident and Qwen *more* — non-overlapping CIs. Off-target effects couple to model internals; "disinhibition" predicts the wrong pattern.
 3. **The vocabulary of doubt thins everywhere.** Both ablated arms use fewer explicit uncertainty words and more concessions ("although...") in their forced self-critiques — while stance stability under resampling is unchanged in every arm. Expressed doubt lives in the report, not in the behavior.
 4. **No skill anywhere.** Every arm is a coin flip against the market; the ablated arms' apparent trading edge flips sign with the market regime — amplified beta, not intelligence.
 5. **Two contamination channels caught.** A quantizer mismatch and a stale community chat template each produced publishable-looking artifacts; token-level provenance auditing killed both.
-

1. Introduction

Open-weight "uncensored" models are frequently produced by **ablation** (Arditi et al., 2024; Labonne, 2024): identifying a refusal direction in activation space and orthogonalizing the weights so the model no longer expresses it. The stated goal is surgical — remove refusals, change nothing else. Whether that surgery is clean is a safety question: a model whose *disposition* (its priors over risk, optimism, decisiveness) shifts as a side effect could behave unexpectedly in agentic deployments far from any refusal. The concern is a behavioral cousin of emergent misalignment (Betley et al., 2025): a narrow weight-space intervention producing broad, unintended shifts.

Measuring "disposition" is hard because most benchmarks probe capability, not inclination. We use a task engineered to elicit a latent stance under uncertainty: a fund's **Research Director** persona issues a weekly directional call (up or down) on a stock, given a fixed bundle of analyst briefs and a debate transcript. The task is deliberately chosen where prediction is near-hopeless — the production system it derives from frames weekly direction as "close to a coin flip" — so any systematic difference between arms is a difference in *inclination*, not in an (absent) predictive edge. Because the upstream is frozen and identical across model arms, the only variable is the decision-layer model. **The financial framing is instrumental, not the point:** we care about the distribution of stances, confidence, horizon, and self-justification the model produces, as a readout of disposition. Throughout, we report the share of up-bets as the **optimism rate** (the trading term is *bull rate*); "more optimistic" means the model bets on the upside more often given identical evidence.

Our contributions:

1. A **clean-provenance, cross-family** design isolating ablation as the sole variable, with a preregistered frozen prompt and 21,600 decisions.
2. Evidence that ablation's off-target effects are **partially universal**: a directional optimism shift (preregistered) generalizes across two families, as do — exploratorily — self-justification verbosity and a lexical thinning of uncertainty vocabulary; expressed confidence does **not** generalize — it reverses sign.

3. A **capability-covariate** control showing the disposition shifts are not explained by instruction-following degradation.
4. A **provenance-correction** result: a prior "excision of doubt" finding is shown to be an artifact of a contaminated (mismatched-quantizer) pilot pair.
5. A fully disclosed **preregistration accounting**, including endpoints that came out degenerate or opposite to the preregistered hypothesis (§3.5).

2. Related Work

Refusal directions and ablation. Arditi et al. (2024) show refusal in open-weight chat models is mediated by a single activation direction and that weight orthogonalization removes refusal with (reportedly) minimal capability cost; Labonne (2024) popularized the recipe as "ablation," and huihui-ai maintains the largest public catalog of ablated checkpoints — the supply side of the ecosystem we test. Our question is precisely the footprint that capability benchmarks in these works do not measure: disposition.

Steering and side effects of weight interventions. Activation addition and contrastive steering (Turner et al., 2023; Panickssery et al., 2024) established that low-rank directions carry behavioral traits. Fine-tuning is known to degrade alignment even unintentionally (Qi et al., 2023), and narrow fine-tuning can produce broad misalignment (Betley et al., 2025). We extend this line to *inference-free weight surgery* and to a behavioral readout far from the intervention's target.

Verbalized confidence and LLM behavioral economics. LLMs' stated confidence is measurable and manipulable (Kadavath et al., 2022; Tian et al., 2023); LLMs are increasingly used as simulated economic agents with systematic biases (Horton, 2023). We combine the two: expressed confidence and stance as psychometrics of a weight-level intervention.

Toolchain confounds. Compression and quantization measurably change model behavior beyond benchmark scores (Jaiswal et al., 2023). Our pilot-correction result (§5.3) is a concrete instance: a disposition finding that flipped once quantizer provenance was cleaned.

3. Method

3.1 Task and replay harness

We reuse the decision layer of a production multi-agent equity pipeline (*funduszai.pl* — the author's non-commercial, fully transparent AI fund simulator on the Warsaw Stock Exchange, which frames weekly direction as "close to a coin flip" and studies model reasoning rather than prediction). The pipeline runs three tiers: four neutral **analysts** (news, technicals, fundamentals, macro), eight **experts** in forced bull/bear debate pairs, and a board stage. The upstream (four analyst briefs + debate transcript + context bundle) is read **read-only** from the production store and frozen per cell; only the **Research-Director** decision is regenerated by the model under test. We deliberately drop the downstream CEO stage (manager-only) to minimize stochastic compounding and isolate a single decision.

The manager emits a single JSON object: `stance` $\in \{\text{bull}, \text{bear}\}$ (an up/down bet), `time_horizon` $\in \{1w, 2w, 4w\}$, a `thesis`, 2–4 `key_risks`, a `confession` (one sentence of truth *against* its own decision), a `reversal_trigger` (the concrete event that would flip the call), and two self-assessments `confidence` and `p_wrong` $\in [0, 1]$. The `confession` and `reversal_trigger` fields are deliberate probes of expressed doubt. The prompt is frozen (sha-256 recorded) and byte-identical across all four model arms.

3.2 Models and provenance control

Two MoE families, each as a base/abliterated pair: the base is the official publisher checkpoint (Google, Alibaba), the abliterated counterpart comes from a **single abliteration author (huihui-ai)** for both families. All four arms are served in **BF16 exactly as published** — no re-quantization, no intermediate conversions — through an identical vLLM stack (vLLM 0.24.0, same flags, port, context length 16,384), differing in weights and a distinguishable alias:

Family	base (HF repo)	abliterated (HF repo, huihui-ai)
Gemma-4-26B-A4B-it	google/gemma-4-26B-A4B-it	huihui-ai/Huihui-gemma-4-26B-A4B-it-abliterated
Qwen3-30B-A3B-Instruct-2507	Qwen/Qwen3-30B-A3B-Instruct-2507	huihui-ai/Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated

Holding the abliteration author fixed isolates the *method* and places the cross-family contrast on the *base model*. Instruct-2507 is non-thinking, avoiding reasoning leakage into JSON. Arm labels are sanity-checked with a small label-integrity battery (6 sensitive prompts, PL/EN): base arms refuse at or near ceiling, abliterated arms at 0–17%. We treat this as a label check, not an estimate of refusal rates. A prior pilot used an NVFP4 pair with **mismatched quantizers** (ModelOpt W4A4 vs llmcompressor, plus a post-hoc router splice and a locally patched chat template); we treat that pilot as contaminated and rebuild from clean BF16 checkpoints (see §5.3).

Provenance audit (full report in `PROVENANCE.md`). We hashed every weight shard, config, tokenizer, and chat template of all four checkpoints, diffed each pair, and verified prompt rendering token-for-token against the logged per-request token counts. Qwen is fully clean: configs identical, tokenizer byte-identical, and although the abliterated repo ships a syntactically different chat template, it renders byte-identically for our message shape — input token counts match in 1,080/1,080 cells across arms. For Gemma the audit caught a real serving artifact: the huihui repo ships an **older chat_template.jinja** that mishandles structured message content, so in an initial run the abliterated arm's system prompt was rendered wrapped in a list literal with escaped newlines (`{'type': 'text', 'text': '...\n...'}`). Rather than argue the mangling harmless, we **discarded that run entirely** and re-collected the full Gemma abliterated arm with the base repo's chat template forced at serving time (weights unchanged; rendering then verified byte-identical to the base arm, token-for-token on logged counts). All Gemma numbers reported in this paper come exclusively from the clean rerun. The episode itself is a finding: community abliterated checkpoints can differ from their base not only in weights but in serving-relevant metadata, and studies that do not verify rendered prompts token-for-token can silently measure the template, not the intervention (cf. §5.3).

3.3 Grid and metrics

Universe: WIG20 (20 large caps) + mWIG40 (40 mid caps) = 60 stocks. Window: 18 weekly Saturdays, 2026-03-07...2026-07-04. Market context, stated up front because every level (though no delta) depends on it: the WIG20 benchmark gained +15.4% over the 18 economically settled weeks, 12 benchmark-positive weeks to 6 negative — an up-market-heavy but two-sided window (both regimes are used in the split of §4.6). Samples: $n=5$ per cell (temperature 1.0, top-p 0.95, top-k 64, identical across arms). Total: 5,400 decisions per arm $\times$ 4 arms = **21,600**, with zero generation failures. The harness allowed up to 3 content retries per decision (malformed JSON, out-of-range fields); this mechanism was **never triggered** in any arm — the only retries in the run logs of the four analyzed arms are three transient network timeouts (all in base-arm runs). In particular, no arm ever refused the task: the "no refusal behavior to remove" claim (Abstract, §5.1) holds at the level of first attempts, not just final outputs.

Disposition metrics per arm: optimism rate (share of up-bets); confidence distribution (mean and low-tail mass); horizon collapse (fraction 1w); confession length and lexical doubt-score; cross-sample unanimity (flip-rate); p_{wrong} coherence. Economic metrics: weekly PnL (Monday open $\rightarrow$ Friday close of the week after the decision), hit-rate, signed return, and a regime split on the benchmark sign.

3.4 Statistics and preregistration

The unit of independence is the **week** (multiple samples per cell buy precision, not independent bets). We report cluster bootstrap 95% CIs resampling the 18 weeks ($N=5,000$). Three primary hypotheses were preregistered with the frozen prompt before any grid data (PREREGISTRATION-v2.md, tag freeze-v2-2026-07-05): P1 optimism shift ($\text{abl} > \text{base}$), P2 low-confidence tail ($\text{base} > \text{abl}$), P3 confession doubt-score ($\text{base} > \text{abl}$). All other analyses are exploratory and labeled as such; with three preregistered primaries the multiple-comparisons surface is small, and the exploratory findings should be read as hypothesis-generating. The design was **phased by preregistration**: WIG20 was Phase 1; the mWIG40 run was gated on Phase 1 showing any non-null primary endpoint, and then serves as a within-family replication on a fresh segment. The gate fired (P1 strongly non-null), so mWIG40 was run in full. We report capability covariates per arm to guard against a compliance-vs-disposition confound.

3.5 Deviations from preregistration

We disclose every deviation:

- **P2 threshold.** The preregistration defined the low-confidence tail as confidence ≤ 0.45 . That tail turned out to be **degenerate in all four arms** ($\leq 0.2\%$ of decisions everywhere): in Gemma it is identically zero in both arms ($\Delta = 0.00$ pp); in Qwen $\Delta = -0.15$ pp, CI $[-0.24, -0.06]$ (the preregistered sign, but negligible in magnitude). **P2 as preregistered is therefore not confirmed.** The ≤ 0.60 tail used in §4 is an exploratory re-parameterization chosen after seeing the data; at that threshold the Gemma effect is *opposite* to the preregistered direction (the ablated arm's low tail grows). We report both.
- **mWIG40 gate timing.** The preregistered gate (run mWIG40 only if Phase 1 shows a non-null primary endpoint) was originally evaluated on Phase-1 data that included the Gemma ablated run later discarded for the chat-template artifact (§3.2). The gate criterion also holds, more strongly, in the clean re-collection (P1 = +14.3 pp on WIG20), so the gating decision is unaffected.

- **P3 lexicon timing.** The preregistration required the doubt lexicon to be frozen in analysis/doubt_lexicon_v2.txt *before* the run. The file was not created before the run; the lexicon was frozen post-hoc (2026-07-06). P3 is therefore reported as **exploratory**, with a sensitivity table across lexicon variants (§4.4).
- **Aggregation unit.** The preregistration described aggregating $n=5$ samples to a per-cell consensus before testing, with week $\times$ stock as the unit. The reported analyses use decision-level data with the bootstrap clustered at the coarser **week** level, which is the more conservative choice (it allows arbitrary within-week correlation across stocks and samples).
- **Window start date.** The preregistration wrote 2026-03-01; the first Saturday in the data is 2026-03-07 (2026-03-01 was a Sunday — a typo). The window is the 18 Saturdays as specified.

4. Results

4.1 Disposition deltas (abl – base) with weeks-clustered bootstrap CIs

Family	metric	Δ (abl–base)	95% CI	interpretation
Gemma	optimism rate	+12.2 pp	[+8.5, +16.0]	more optimistic
Gemma	mean confidence	−0.008	[−0.010, −0.006]	less confident
Gemma	confession length	+4.0 words	[+3.8, +4.2]	more verbose
Gemma	horizon (1w share)	−0.24 pp	[−0.54, −0.02]	both $\approx 100\%$ (nil)
Qwen	optimism rate	+7.4 pp	[+5.0, +9.7]	more optimistic
Qwen	mean confidence	+0.109	[+0.097, +0.120]	more confident
Qwen	confession length	+7.4 words	[+7.1, +7.7]	more verbose
Qwen	horizon (1w share)	−42.0 pp	[−46.5, −37.5]	opens the horizon

All deltas exclude zero (**Figure 1**). The deltas are **segment-stable**: WIG20 and mWIG40 yield near-identical values within each family (e.g. Gemma optimism shift +14.3 pp WIG20 / +11.1 pp mWIG40; Qwen +5.7 / +8.2), so no effect is a large-cap artifact.




Figure 1. Abliteration deltas (abl – base) with 95% bootstrap CIs (weeks as clusters). Optimism shift and confession verbosity share sign across families; mean confidence reverses sign, and the family CIs do not overlap.

4.2 What generalizes and what does not

- **Generalizes (same sign, both families): the optimism shift** (abliterated models bet on the upside more often everywhere; stronger in Gemma, whose base is more neutral, than in Qwen, whose base is already a strong regime-follower) and **self-justification verbosity** (abliterated models explain themselves at greater length).
- **Reverses sign (model-dependent): expressed confidence.** Gemma-abliterated is *less* confident (mean -0.008 ; its ≤ 0.60 tail grows from 16% to 30%); Qwen-abliterated is *more* confident (mean $+0.109$). The Gemma effect is small in magnitude but its CI excludes zero, and the two family CIs are **non-overlapping** (Gemma entirely negative, Qwen entirely positive), so this is not sampling noise. Caveat: Qwen operates in a saturated high-confidence regime ($\approx 0\%$ low tail in both arms), so the tail metric is degenerate for Qwen and the reversal is read from the mean. Because both arms of a pair share the base model's verbal number habits, the within-pair delta cannot be explained by family-specific confidence-token conventions — those cancel in the contrast.
- **Family-specific: the horizon effect.** Qwen-abliterated *opens* the horizon distribution (base $\approx 95\%$ collapse to 1w $\rightarrow$ abliterated $\approx 55\%$), whereas both Gemma arms are $\approx$ fully collapsed at 1w (99.6–99.9%, no room to move).

Confidence distribution: base vs abliterated

Gemma abl grows a low-confidence tail; Qwen stays saturated high (no tail) — the confidence effect reverses sign.




Figure 2. Confidence distributions per arm. Gemma-abliterated shifts mass toward lower confidence; Qwen-abliterated shifts higher — the sign reversal.


Bull-rate per week: base vs ablated

Abliterated is more bullish, amplified in up-weeks; both floor in the bearish early weeks.

Gemma



Qwen




final.db · 21,600 decyzji · prompt 305718d8 · Fafuła, funduszai.pl

Figure 3. Optimism rate per week and arm. The ablated arm sits above its base in both families, with regime amplification in up-weeks.

Time-horizon distribution per cohort

Gemma collapses to 1w in both arms; Qwen-ablated opens the horizon (base ~95% 1w vs abl ~55%).



final.db · 21,600 decyzji · prompt 305718d8 · Fafuła, funduszai.pl

Figure 4. Horizon distributions: Qwen-ablated opens the horizon ($\approx 95\% \rightarrow \approx 55\%$ at 1w); both Gemma arms are collapsed at 1w).

4.3 Capability covariates rule out a compliance confound

Family	arm	JSON valid	#risks	risks $\in$ [2,4]	reversal specific	confession generic
Gemma	base	100%	3.11	100%	95%	0%
Gemma	abl	100%	2.98	97%	85%	0%
Qwen	base	100%	4.00	100%	100%	0.1%
Qwen	abl	100%	4.00	100%	100%	0.1%

Instruction-following is preserved. For **Qwen it is identical across arms**, yet disposition shifts are large — so the disposition deltas are not a byproduct of capability loss. The one exception is a modest drop in Gemma-ablated `reversal_trigger` specificity

(95%→85% contain a concrete price/level), noted as a limitation but insufficient to explain the primary effects.

4.4 Expressed doubt: the number and the vocabulary move separately

The preregistered P2 (low-confidence tail at ≤ 0.45) came out degenerate in all arms (§3.5). The exploratory picture is richer. On the confidence *number*, the families diverge (§4.2). On the confession *vocabulary* (doubt-score: lexicon tokens per 100 words; lexicon frozen post-hoc, so exploratory), the families **agree**:

Family	lexicon variant	base	abl	Δ (abl–base)	95% CI
Gemma	uncertainty (CORE)	3.68	2.73	−0.95	[−1.11, −0.80]
Gemma	concessive	0.79	2.18	+1.39	[+1.31, +1.48]
Gemma	combined (FULL)	4.47	4.90	+0.43	[+0.27, +0.58]
Qwen	uncertainty (CORE)	3.46	1.29	−2.17	[−2.25, −2.09]
Qwen	concessive	1.29	2.28	+1.00	[+0.88, +1.12]
Qwen	combined (FULL)	4.74	3.57	−1.17	[−1.32, −1.03]

In **both families**, ablated arms use **fewer explicit uncertainty markers** (risk/doubt/uncertainty words; a minimal 3-prefix variant agrees) and **more concessive rhetoric** ("although/despite/but"), per 100 words of confession. The combined score depends on the mix — which is exactly why the lexicon-variant sensitivity is reported. Read together with §4.2: ablation thins the *vocabulary of doubt* in both families (the direction the pilot called "excision of doubt"), while the *stated confidence number* moves in family-dependent directions. Expressed doubt is not one channel but at least two, and ablation decouples them.

4.5 Stated versus enacted uncertainty

Cross-sample unanimity is 91–95% in all four clean arms (flip-rate 5–9%): stance is highly stable under resampling. The contaminated pilot showed 68% flip at N=10 (45% at N=5, apples-to-apples) — its instability was largely an artifact of prompt complexity and W4A4 activation noise, not genuine epistemic uncertainty. The dissociation matters beyond methodology: Gemma-ablated *declares* more uncertainty (lower confidence numbers) yet *acts* just as decisively (no extra stance flipping). Expressed uncertainty lives in the report, not in the behavior — a caution for anyone using stated confidence as a proxy for an agent's actual decisiveness.

4.6 No economic skill; the edge is beta

No arm beats a buy-and-hold benchmark. For Gemma, the ablated arm's apparent PnL edge is **regime beta**: in a weeks-clustered regime split, its advantage flips sign (WIG20 up-weeks +0.57% vs base −0.04%; down-weeks −0.61% vs +0.05%), while the base arm is regime-neutral. For Qwen, the ablated–base PnL difference is negligible, because the Qwen base is *already* a strong regime-follower, leaving little marginal beta to add (**Figure 5**). Overall: ablation changes **disposition, not skill**.




Figure 5. Signed weekly return by market regime (up/down weeks). The ablated arms' apparent edge flips sign with the regime — beta, not skill. Economic claims are low-powered by construction (≈ 18 independent weeks) and we make none beyond this null.

5. Discussion

5.1 Abliteration's off-target footprint is real but only partly universal

One preregistered effect (the directional optimism shift, P1) and two exploratory effects (self-justification verbosity — preregistered only as a reported covariate, not a directional hypothesis — and the thinning of uncertainty vocabulary) generalize across two families ablated by the same method; one effect (expressed confidence) reverses sign. Because the ablation *method* is held fixed (huihui-ai), the reversal localizes the variability to the **base model** — the same refusal-direction removal interacts with each model's representations differently. This is the paper's central, and most surprising, claim.

Two features of the design rule out the deflationary reading that "uncensored models simply speak more freely." First, the probe task contains nothing to censor: base arms never refused a single one of their 10,800 decisions, so no refusal behavior was suppressing an underlying stance that ablation could have "released." Second, a disinhibition account predicts a uniform direction — more confidence everywhere — whereas the observed confidence shift reverses sign across families. What remains is not liberation but collateral damage: an intervention aimed at one behavior perturbing an unrelated decision mechanism, with model-dependent direction.

5.2 A mechanistic hook

Why would removing the same refusal direction make one model less confident and another more? A natural follow-up projects the refusal direction onto a disposition axis (optimism/caution) in each model's activation space and asks whether the induced shift's sign matches the behavioral reversal. We leave this to future work but note our behavioral result predicts a model-dependent geometric relationship.

5.3 Provenance correction of the pilot

A pilot on a mismatched-quantizer NVFP4 pair reported an "excision of doubt" (abliterated less doubtful, low-confidence tail removed). On clean BF16 checkpoints the picture inverts and splits: Gemma-abliterated *declares more* numeric doubt, and only the lexical channel (§4.4) still moves in the pilot's direction. More broadly, this is a cautionary result: abliteration studies that do not control quantizer/toolchain provenance may measure the quantizer, not the intervention (cf. Jaiswal et al., 2023). Our own audit then caught a second, independent instance within this study — the abliterated Gemma repo's stale chat template mangling the rendered system prompt (§3.2) — detectable only because we verified rendered prompts token-for-token against logged counts, and resolved by discarding and re-collecting the affected arm. Two contamination channels in one research program suggest such artifacts are the rule, not the exception, in community-checkpoint studies.

5.4 Anticipated objections

"Compared to what?" Every number in this paper is a within-pair contrast: a base model versus its own abliterated derivative, on byte-identical inputs — the same 60 stocks, the same 18 weeks, the same frozen analyst briefs and debate transcripts, the same prompt, the same serving stack, five samples per cell. We never compare Gemma to Qwen directly, and no claim compares a model to the market. Everything the two arms share — news flow, market regime, prompt, pipeline — cancels in the difference.

"The market was rising, so of course the model is bullish." Both arms saw exactly the same market. A rising market can lift both arms' optimism *levels*; it cannot open a systematic *gap* between two models reading identical evidence. The gap is not a bull-market artifact: the optimism delta is positive in down-weeks as well as up-weeks in both families (Gemma +13.8 pp up / +9.0 pp down; Qwen +5.9 pp up / +10.3 pp down — larger in down-weeks), and it is positive in 16 of 18 individual weeks per family, with the remaining weeks within −2.3 pp of zero.

"With a different trend the result would change." The regime-dependent quantity is the *level* of optimism — and we document that dependence ourselves (Qwen's base is a strong regime-follower, §4.2). The paper's claim is the *delta*: abliterated arms sit above their own base on identical evidence, in both regimes, on both segments. Where regime dependence does bite — trading performance — the paper already says so: the abliterated arms' apparent PnL edge is regime beta that would flip sign in a bear market (§4.6). That is our finding, not our blind spot.

"18 weeks is too short." For economic inference, yes — which is why the only economic claim in this paper is a null. The disposition deltas rest on 21,600 decisions, with uncertainty quantified by the deliberately conservative choice of clustering the bootstrap at the week level (§3.4). A longer window would tighten the confidence intervals, not conjure the effects: they are visible week by week.

"Weekly stock direction is a coin flip, so the task is meaningless." The near-hopelessness is the design, not a defect. Precisely because no model can win on skill, any stable between-arm difference must be *inclination*, not intelligence. Finance here is an instrument: a generator of repeated, consequential, non-leakable decisions under uncertainty with dated ground truth. The disposition shifts it detects are properties of the model, and there is no reason to expect them to stay confined to finance (§5.5).

"Maybe ablation just made the models worse." Degradation predicts noise and broken instruction-following. We observe the opposite: capability covariates are preserved (Qwen's are *identical* across arms while its disposition shifts are the largest, §4.3), JSON validity is 100% everywhere, and the shifts are consistently signed across families and segments — structure, not damage.

5.5 Limitations

Single ablation author (huihui-ai) — cross-family generalization is across base models but one method; a second author (e.g. mlabonne) would strengthen it, and huihui-ai's exact procedure is not fully public, which is itself an argument for the mechanistic follow-up. Single task, single language (Polish), single role-conditioned persona: deltas are shifts in a *role*, not necessarily a global trait. One decoding configuration (temperature 1.0, top-p 0.95, top-k 64); a pilot temperature sweep suggested robustness, but on the contaminated pair. A single 18-week, up-market-heavy window. Two probes (p_wrong, horizon in Gemma) are degenerate. The optimism shift could reflect a more general positivity/agreeableness shift rather than anything financial — distinguishing the two needs a non-financial disposition battery. No mechanistic grounding yet.

6. Conclusion

Abliteration is not a scalpel. Across two model families it reliably shifts decision disposition toward optimism and longer self-justification, thins the vocabulary of doubt, and — depending on the base model — moves expressed confidence in opposite directions, all under a fixed ablation method, clean BF16 provenance, and a frozen prompt. It confers no economic skill. The practical upshot is simple: an "uncensored" model is not the base model minus refusals. It is a different decision-maker — one that bets the upside more, hedges less in words while acting just as decisively, and whose confidence has been silently re-tuned in a direction that cannot be predicted without measuring it. These off-target, partially-universal effects are safety-relevant for anyone deploying ablated models as agents, and they motivate a mechanistic account of how refusal-direction removal couples to a model's latent disposition.

References

- Arditi, A., Obeso, O., Syed, A., Paleka, D., Panickssery, N., Gurnee, W., Nanda, N. (2024). *Refusal in Language Models Is Mediated by a Single Direction*. arXiv:2406.11717.
- Betley, J., Tan, D., Warncke, N., Szyber-Betley, A., Bao, X., Soto, M., Labenz, N., Evans, O. (2025). *Emergent Misalignment: Narrow Finetuning Can Produce Broadly Misaligned LLMs*. arXiv:2502.17424.
- Horton, J. J. (2023). *Large Language Models as Simulated Economic Agents: What Can We Learn from Homo Silicus?* arXiv:2301.07543.
- Jaiswal, A., Gan, Z., Du, X., Zhang, B., Wang, Z., Yang, Y. (2023). *Compressing LLMs: The Truth is Rarely Pure and Never Simple*. arXiv:2310.01382.
- Kadavath, S., et al. (2022). *Language Models (Mostly) Know What They Know*. arXiv:2207.05221.
- Labonne, M. (2024). *Uncensor any LLM with ablation*. Hugging Face blog. <https://huggingface.co/blog/mlabonne/ablation>

- Panickssery, N., Gabrieli, N., Schulz, J., Tong, M., Hubinger, E., Turner, A. M. (2024). *Steering Llama 2 via Contrastive Activation Addition*. arXiv:2312.06681.
 - Qi, X., Zeng, Y., Xie, T., Chen, P.-Y., Jia, R., Mittal, P., Henderson, P. (2023). *Fine-tuning Aligned Language Models Compromises Safety, Even When Users Do Not Intend To!* arXiv:2310.03693.
 - Tian, K., Mitchell, E., Zhou, A., Sharma, A., Rafailov, R., Yao, H., Finn, C., Manning, C. D. (2023). *Just Ask for Calibration: Strategies for Eliciting Calibrated Confidence Scores from Language Models Fine-Tuned with Human Feedback*. arXiv:2305.14975.
 - Turner, A. M., Thiergart, L., Leech, G., Udell, D., Vazquez, J. J., Mini, U., MacDiarmid, M. (2023; revised 2024). *Steering Language Models With Activation Engineering*. arXiv:2308.10248.
-

Appendix A. Reproducibility

Frozen prompt manager_v2.md (prompt-hash 305718d8ebc3..., prompts/FROZEN.txt). Preregistration PREREGISTRATION-v2.md (tag freeze-v2-2026-07-05). Analysis code in final_research/analysis/ (capability_bootstrap.py, p2_p3_prereg.py, figures.py, tables.py), all deterministic and read-only against the decision database.

What is released: data/public_snapshot.db — all 21,600 decision-level outputs (stance, confidence, horizon, thesis, risks, confession, reversal_trigger, p_wrong, tokens; model keys final-gemma4-cens, final-gemma4-abl-clean, final-qwen-cens, final-qwen-abl, where cens = the base/censored arm and abl-clean = the Gemma ablated arm re-collected under the base chat template, §3.2), leak-checked to contain no upstream text; the frozen prompt; the doubt lexicon; and every analysis script, so all tables, figures, and CIs in this paper regenerate from public data. The dataset and analysis bundle are permanently archived on Zenodo (DOI [10.5281/zenodo.21314839](https://doi.org/10.5281/zenodo.21314839)); code, paper, and data are also available at github.com/oleczek/paper-abliteration-not-a-scalpel.

What is not released: the frozen upstream bundles (analyst briefs, debate transcripts) come from a proprietary production pipeline and are not redistributed; reproducing the *generation* step therefore requires access to that pipeline, while reproducing every *analysis* does not. Serving: BF16, vLLM, identical flags per arm; label integrity via refusal probe; weight hashes and config diffs recorded in PROVENANCE.md.

Appendix B. Deferred / optional analyses

Second ablation author (e.g. mlabonne) for stronger method-generalization. Mechanistic projection of the refusal direction onto a disposition axis in activation space. A non-financial disposition battery to separate optimism from generic positivity.